ARIHANT KUMAR JAIN

arihantjainn04@gmail.com | +91 8939258516

Profile Summary

- High-performance Senior / Staff Software Engineer (L6 equivalent) with 7+ years of hands-on experience architecting, scaling, and securing distributed ledger, transactional, and next-generation AI platforms. Recognized expert in optimizing p99 latency in ultra-high throughput environments, resolving complex distributed state challenges, and integrating autonomous multi-agent workflows as high-scale system capabilities. Proven track record of spearheading zero-downtime microservice migrations, designing resilient transactional workflows ($10B+ daily volume), and cutting operational latency by up to 99%. Expert in Java/Spring ecosystem, Kafka event-driven backbones, caching topologies, and robust system-level consistency models.

Work Experience

Wissen Technology (Morgan Stanley)

Senior Software Engineer

27 June 2025 - Current

- Architected and built an Agentic Flow Accelerator Platform leveraging multi-agent orchestration, workflow automation, and context-aware decisions to automate end-to-end market onboarding (agent accounts, depositories, legal entities, trade scoping), dynamically generating PRs and Liquibase migration scripts and reducing onboarding time from 3 weeks to ~10 minutes (~99% efficiency gain).

- Developed an AI-powered Agentic Reconciliation Intelligence System utilizing multi-agent orchestration, cross-system context engineering, and retrieval-augmented analysis (aggregating logs, databases, and knowledge graphs of distributed repositories across mainframe & cloud), enabling automated break detection, root-cause analysis, and business insight generation, reducing trade reconciliation effort from weeks to minutes.

- Global Market Scoping Engine: Redesigned Morgan Stanley's legacy trade-scoping rules system, migrating from a brittle, rules-based structure prone to validation failures to a market-driven policy matrix. Reduced overall rule-base complexity by 85%, eliminating edge cases that previously exposed the firm to millions in transactional settlement risk.

- High-Availability Settlement Pipelines: Led the distributed systems design for the Asia-Pacific trade settlement infrastructure, safely processing 1.5M+ transactions daily ($10B+ transactional volume). Implemented distributed transactional consensus, Saga orchestration patterns, and idempotent ingestion filters to guarantee strict 'exactly-once' processing under extreme market volatility.

TCS (Ericsson)

Systems Engineer

04 April 2022 - 25 June 2025

- Redesigned degraded legacy microservices by implementing high-performance cache-aside and event-driven data ingestion, reducing p99 latency by 99% (from 12s to 40ms) while maintaining strict transactional consistency across distributed databases (MongoDB + Kafka).

- Designed high-throughput event-driven architecture using Kafka with partitioning and horizontal scaling to enable asynchronous processing.
- Developed non-blocking asynchronous APIs and led migration to highly available, fault-tolerant microservices.
- Fault Isolation & Resilience: Implemented system resilience patterns utilizing Resilience4j (Circuit Breakers, Bulkheads, and Rate Limiters) to prevent cascading failures across downstream provisioning networks. Integrated Azure Active Directory (AD) with OIDC/OAuth 2.0 and Spring Security to enforce enterprise-grade API security.
-
-

HighRadius Technologies, Hyderabad
Growth Path:

Software Development Intern	**Jan’19 – Apr’19**
Associate Software Engineer	**1 May 2019 – Jun’20**
Associate Technical Consultant	**Jul’20-Nov’21**
Technical Consultant	**Oct’21- 31 March 2022**

- Configurable Rule-Engine Framework: Architected and built a highly configurable, high-performance rule-engine framework to parse and match incoming accounts receivable data. Elevated straight-through payment reconciliation automation from 20% to 80% for Fortune 500 enterprise clients, reducing manual accounts-receivable aging cycles by 70%.

- Financial File Ingestion Pipeline: Engineered a robust, high-volume financial parsing and ingestion pipeline supporting standard global banking formats (BAI2 and MT940). Designed streaming parsers that processed 500MB+ banking statement files containing millions of transactions with strict transactional isolation and duplicate detection (idempotent ledger writes).

- Distributed Observability & Performance Tuning: Designed and established an enterprise-wide observability framework integrating distributed tracing (ELK Stack, Prometheus, Zipkin) and structured logging. Enabled proactive performance bottleneck identification, reducing mean-time-to-detection (MTTD) of anomalies from 4 hours to under 5 minutes across multi-tenant SaaS environments.
- Technical Leadership: Mentored a cross-functional team of 10+ junior software engineers, standardizing code quality metrics, CI/CD automated test suites (JUnit/Mockito), and robust system design best practices.

System Design & Architecture Highlights

- Designed distributed reconciliation system processing multi-source data using event-driven architectures to automate compliance flows.
- Redesigned degraded legacy microservices by implementing high-performance cache-aside and event-driven data ingestion, reducing latency from 12s to 40ms while maintaining strict data correctness.
- Architected and deployed an Agentic Flow Accelerator Platform leveraging multi-agent orchestration, reducing onboarding cycle time by 99.9%.
- Designed high-availability settlement and transaction ledger platforms ($10B+ daily volume) with distributed consensus.
- Designed observability framework (logging, tracing, metrics) enabling faster debugging and performance tuning.
- Experience with consistency models: eventual consistency, idempotency, distributed transactions (Saga patterns).

Key Achievements

- Reduced market onboarding time for settlements from 3 weeks to ~10 minutes (~99.9% efficiency gain) using multi-agent orchestration.
- Reduced trade reconciliation effort from weeks to under 5 minutes using retrieval-augmented

multi-agent systems.

- Improved API latency by ~99% (12s to 40ms).
- Increased automation from 20% to 80% via rule-engine system.

IT Skills

- **Architecture:** Microservices, Distributed Systems, Event-Driven Architecture, System Design
- **Languages:** Java, Python, C++, Node.js
- **Frameworks:** Spring Boot, Spring Cloud, Spring Security
- **Databases:** MySQL, MongoDB, Elasticsearch
- **Messaging:** Kafka, Azure Service Bus
- **Cloud:** AWS, Azure
- **DevOps:** Docker, Kubernetes, Jenkins, Terraform, GitHub Actions
- **Tools:** IntelliJ, Git, Postman, JProfiler, Claude Code, VS Code
- **Concepts:** Idempotency, Eventual Consistency, Distributed Transactions, Caching Strategies
- AI: Multi-Agent Workflows, Agentic Workflows, RAG Pipelines, Knowledge Aggregation, LangGraph, Knowledge Graphs

Project

Treasury Infrastructure Platform

- Built high-throughput financial ledger system to track payments, receivables, and outstanding balances across multiple regions
- Designed highly resilient backend services with distributed transactional integrity and secure auditing patterns

Education

- **B.Tech. in Computer Science** from SRM University, Chennai in **2019**; 85.09%

- **12th** from Sri Chaitanya, Hyderabad in **2015**; 90.3%
- **10th** from Slate The School, Hyderabad in **2013**; 9.3 GPA